

Paxos – Sample Input 2

Suppose we want to simulate a scenario where there are two proposers, one of which fails and later recovers. We define the simulation as follows:

$$(n_P = 2, n_A = 3, t_{\max} = 50, E)$$

Where E contains four events:

- $(t = 0, F = \emptyset, R = \emptyset, \pi_c = p_1, \pi_v = 42)$
- $(t = 8, F = \{p_1\}, R = \emptyset, \pi_c = \emptyset, \pi_v = \emptyset)$
- $(t = 11, F = \emptyset, R = \emptyset, \pi_c = p_2, \pi_v = 37)$
- $(t = 26, F = \emptyset, R = \{p_1\}, \pi_c = \emptyset, \pi_v = \emptyset)$

In other words, we have a system with two Proposers and three Acceptors. p_1 proposes value 42 at tick 0, but fails at tick 8. Then, at tick 11, p_2 proposes a different value (37). Finally, p_1 recovers at tick 26.

Ticks 0–7

Same as in Example 1.

Contents of N at the end of tick 7			
$m.type$	$m.src$	$m.dst$	Other attributes
ACCEPT	p_1	a_2	proposal_id=1, value=42
ACCEPT	p_1	a_3	proposal_id=1, value=42
ACCEPTED	a_1	p_1	proposal_id=1, value=42

Tick 8

In this tick, there is an event which specifies that p_1 will go into a failed state (remember that this happens *before* checking the contents of the network). So, since all the messages in N have p_1 as either the source or the destination, they cannot be delivered.

Contents of N at the end of tick 8			
$m.type$	$m.src$	$m.dst$	Other attributes
ACCEPT	p_1	a_2	proposal_id=1, value=42
ACCEPT	p_1	a_3	proposal_id=1, value=42
ACCEPTED	a_1	p_1	proposal_id=1, value=42

Ticks 9 and 10

During these ticks, no messages are delivered. However, the simulation carries on because t_{\max} hasn't been reached, and there are still messages pending delivery.

Ticks 11-14

In tick 11, there is an event specifying that p_2 must propose the a new value (37). Ticks 11-14 are similar to ticks 0-3, except p_2 will use a larger proposal number (2). Additionally, a_1 had previously accepted a value (value 42 in proposal 1). So, when sending a PROMISE message to p_2 , it specifies the prior value it accepted (and its proposal number).

Contents of N at the end of tick 14			
$m.type$	$m.src$	$m.dst$	Other attributes
ACCEPT	p_1	a_2	proposal_id=1, value=42
ACCEPT	p_1	a_3	proposal_id=1, value=42
ACCEPTED	a_1	p_1	proposal_id=1, value=42
PROMISE	a_1	p_2	proposal_id=2, prior_proposal=($v = 42, p = 1$)
PROMISE	a_2	p_2	proposal_id=2, prior_proposal= \emptyset
PROMISE	a_3	p_2	proposal_id=2, prior_proposal= \emptyset

Ticks 15-16

Similar to ticks 4-5, except p_2 cannot use the value specified in the PROPOSE message. Since one of the Acceptors had previously accepted value 42, then p_2 is bound to use that value.

Contents of N at the end of tick 16			
$m.type$	$m.src$	$m.dst$	Other attributes
ACCEPT	p_1	a_2	proposal_id=1, value=42
ACCEPT	p_1	a_3	proposal_id=1, value=42
ACCEPTED	a_1	p_1	proposal_id=1, value=42
ACCEPT	p_2	a_1	proposal_id=1, value=42
ACCEPT	p_2	a_2	proposal_id=1, value=42
ACCEPT	p_2	a_3	proposal_id=1, value=42

Tick 17-23

Similar to ticks 6-12 in Example 1. At the end of these ticks, p_2 has determined there is consensus around value 42.

Contents of N at the end of tick 23			
$m.type$	$m.src$	$m.dst$	Other attributes
ACCEPT	p_1	a_2	proposal_id=1, value=42
ACCEPT	p_1	a_3	proposal_id=1, value=42
ACCEPTED	a_1	p_1	proposal_id=1, value=42

Ticks 24-25

During these ticks, no messages are delivered. However, the simulation carries on because t_{\max} hasn't been reached, and there are still messages pending delivery in N and future events in E .

Tick 26

In this tick, there is an event specifying that p_1 should recover. This means we can deliver the **ACCEPT** message from p_1 to a_2 . However, at this point, a_2 has already promised to not participate in any proposals with a number less than 2. So, it sends a **REJECTED** message back to p_1 .

Contents of N at the end of tick 26

$m.type$	$m.src$	$m.dst$	Other attributes
ACCEPT	p_1	a_3	proposal_id=1, value=42
ACCEPTED	a_1	p_1	proposal_id=1, value=42
REJECTED	a_1	p_1	proposal_id=1

Tick 27

Same as tick 27, but with Acceptor a_3 .

Contents of N at the end of tick 27

$m.type$	$m.src$	$m.dst$	Other attributes
ACCEPTED	a_1	p_1	proposal_id=1, value=42
REJECTED	a_2	p_1	proposal_id=1
REJECTED	a_3	p_1	proposal_id=1

Tick 28-30

In these ticks, the **ACCEPTED** message from a_1 (which had been delayed since tick 7) and the **REJECTED** messages from a_2 and a_3 are delivered. Confronted with a majority of **REJECTED** messages from the acceptors, p_1 realizes its proposal has failed, and decides to initiate a new proposal for value 42, but using proposal number 3.

Tick 31-42

Similar to ticks 1-12 in Example 1, except that the **PROMISE** messages from the Acceptors will all include the previously accepted proposal (2) and value (42). The algorithm still runs through the **ACCEPT/ACCEPTED** phase, but is guaranteed to produce consensus around the same value that p_2 reached consensus on.